

MECH&AE 298 Mini-Project 2 Report

William Hsieh

February 16, 2025

1 Introduction

This mini-project infers the result of championship games between two top teams. It does so by analyzing regular season games, and then simulating 500 two-leg finals. A Bayesian statistical model is used in order to calculate the probabilities of each team winning, losing, or drawing a match. A multi-level model is used, indicating a hierarchical structure of parameters where teams "share" statistics.

2 Our Model

2.1 Multi-Level Model

First, the hyperparameters are set: the mean attack and defense are sampled from normal distributions, the variance of attack and defense are sampled from an exponential distribution, and the home advantage is sampled from a normal distribution. These statistics are used to determine the attack and defense for each of the 20 teams.

Compared to other methods of pooling, this multi-level model uses partial pooling. A model that uses no pooling would sample each team's attack and defense power independently, instead of using the statistics sampled from our hyperparameters. On the other hand, a model that uses complete pooling would end up with teams that share the same attack and defense power, instead of having slightly different attack and defense powers like we do here.

2.2 Description

The mean attacking power, mean defending power, variance of attacking power, variance of defending power, and home advantage are defined as follows:

$$\begin{aligned}\mu_{\text{att}} &\sim \mathcal{N}(0, 0.1) \\ \mu_{\text{def}} &\sim \mathcal{N}(0, 0.1) \\ \sigma_{\text{att}} &\sim \exp(1) \\ \sigma_{\text{def}} &\sim \exp(1) \\ \text{home} &\sim \mathcal{N}(0, 1)\end{aligned}$$

These are then used to create a distribution to sample from, resulting in each team's attack and defense power:

$$\begin{aligned}\text{att}[\text{team_i}] &\sim \mathcal{N}(\mu_{\text{att}}, \sigma_{\text{att}}) \\ \text{def}[\text{team_i}] &\sim \mathcal{N}(\mu_{\text{def}}, \sigma_{\text{def}})\end{aligned}$$

Finally, in a head to head, each team's attack and defense are used to calculate the win probabilities for each team:

$$\text{win home} = \text{home advantage} + \text{attack}[\text{home}] + \text{defense}[\text{away}] - \text{offset}$$

$$\text{win away} = \text{attack}[\text{away}] + \text{defense}[\text{home}] - \text{offset}$$

Note that offsets (calculated via $\mu_{\text{att}} + \mu_{\text{def}}$) are added to center the distribution. Finally, the scores can be calculated using the win probabilities:

$$\text{score} \sim \text{Poisson}(\exp(\text{win}))$$

2.3 Championship Games

After 3000 samples of matchups between the 20 teams, the best two teams are selected. Now, we can sample 500 "finals" between the two teams. Team A and team B take turns being the home team; the win probabilities and scores are calculated using the same methods as above.

However, there is one key difference when calculating the scores. Instead of updating the parameters for each team after each game, the parameters remain unchanged. The scores are instead sampled from the Poisson distribution 500 times for each team, to simulate 500 matches without changes in parameters:

$$\text{score} = \text{rand}(\text{Poisson}(\exp(\text{win})), 500)$$

The result is a list of scores, one for each team for the 150,000 (3000×500) championship games. (The 3000 comes from the previous step with the 20 teams.)

3 Results

Plotting the scores of team A (Manchester City) and team B (Manchester United) gives the following:




Figure 1: Heatmap of Scores

Each cell represents the number of points each team scored in a game. For example, the whitest spot is the most common, while the darker and non-populated spots show a low frequency. The

one brightest spot represents a game where team A scored 3 points and team B scored 2 points, resulting in a win for team A.

By counting the number of occurrences where team A scored more than team B, we can figure out the win ratio for team A, and vice versa for team B. If they scored the same amount of points (represented by the diagonal starting at $(0, 0)$, it results in a draw.

- Win probability for team A: 61.504%
- Win probability for team B: 24.090%
- Draw probability: 14.405%

As a sanity check, all of these probabilities add up to 100%.

4 Discussion and Conclusion

Based on the heatmap, we can see that most of the games had around 2-4 points, but in some rare cases could go up to 17 points. The games were also relatively close, with most of the probability mass around the slope $m = 1$. However, this small difference was enough to correspond to a win probability of 61.504% for team A and a much lower win probability of 24.090% for team B.

Overall, this mini-project analyzed games between 20 teams, picked the best 2 teams out from there, and simulated many championship games between these two teams. Bayesian statistics is used to calculate the win probabilities and scores for each game, and a multi-level model was used to convert raw team statistics into a quantifiable score.

5 Code

The Jupyter notebook used is appended to this report.

```
In [15]: using JSON
using DataFrames
using StatsPlots
using Turing
using LinearAlgebra
using Random
```

Multi-level model using football match simulation as an example

```
In [16]: ## First, import the data and do some data wrangling

england_league = JSON.parsefile("../data/matches_England.json")

matches_df = DataFrame{home = [], away = [], score_home = [], score_away = []}
```

0x4 DataFrame

Row	home	away	score_home	score_away
	Any	Any	Any	Any

```
In [17]: # example entry for each game in england_league: "label" => "Burnley
matches = []
for match in england_league
    push!(matches, split(match["label"], ",")) # "Burnley - AFC Bournemouth"
end

for match in matches
    home, away = split(match[1], " - ") # "Burnley" # "AFC Bournemouth"
    score_home, score_away = split(match[2], " - ") # "1" # "2"
    push!(matches_df, [home, away, parse{Int}(score_home), parse{Int}(score_away)])
end

matches_df

teams = unique(collect(matches_df[:,1]))
```

```

20-element Vector{Any}:
"Burnley"
"Crystal Palace"
"Huddersfield Town"
"Liverpool"
"Manchester United"
"Newcastle United"
"Southampton"
"Swansea City"
"Tottenham Hotspur"
"West Ham United"
"Manchester City"
"Leicester City"
"Chelsea"
"Arsenal"
"Everton"
"AFC Bournemouth"
"Watford"
"West Bromwich Albion"
"Stoke City"
"Brighton & Hove Albion"

```

```

In [ ]: ## Now, our model

@model function football_matches(home_teams, away_teams, score_home, s

    # Hyper priors
    μatt ~ Normal(0, 0.1)
    μdef ~ Normal(0, 0.1)
    σatt ~ Exponential(1)
    σdef ~ Exponential(1)
    home ~ Normal(0, 1)

    # Team-specific effects

    att = zeros(length(teams))
    def = zeros(length(teams))

    for i in 1:length(teams)
        att[i] ~ Normal(μatt, σatt)
        def[i] ~ Normal(μdef, σdef)
    end

    # att ~ filldist(Normal(μatt, σatt), length(teams)) # more compact
    # def ~ filldist(Normal(μdef, σdef), length(teams))

    offset = mean(att) + mean(def)

    # the number of matches
    n_matches = length(home_teams)

    # scoring rates θ
    θ_home = Vector{Real}(undef, n_matches)      # or just θ_home = zer

```


Sampling	91%		ETA: 0:00:12
Sampling	92%		ETA: 0:00:12
Sampling	92%		ETA: 0:00:11
Sampling	92%		ETA: 0:00:10
Sampling	93%		ETA: 0:00:10
Sampling	94%		ETA: 0:00:09
Sampling	94%		ETA: 0:00:08
Sampling	94%		ETA: 0:00:08
Sampling	95%		ETA: 0:00:07
Sampling	96%		ETA: 0:00:06
Sampling	96%		ETA: 0:00:05
Sampling	96%		ETA: 0:00:05
Sampling	97%		ETA: 0:00:04
Sampling	98%		ETA: 0:00:03
Sampling	98%		ETA: 0:00:03
Sampling	98%		ETA: 0:00:02
Sampling	99%		ETA: 0:00:01
Sampling	100%		ETA: 0:00:01
Sampling	100%		Time: 0:02:16
Sampling	100%		Time: 0:02:16

Chains MCMC chain (3000×57×1 Array{Float64, 3}):

Iterations = 1001:1:4000
Number of chains = 1
Samples per chain = 3000
Wall duration = 136.29 seconds
Compute duration = 136.29 seconds
parameters = μ att, μ def, σ att, σ def, home, att[1], def[1], att[2], def[2], att[3], def[3], att[4], def[4], att[5], def[5], att[6], def[6], att[7], def[7], att[8], def[8], att[9], def[9], att[10], def[10], att[11], def[11], att[12], def[12], att[13], def[13], att[14], def[14], att[15], def[15], att[16], def[16], att[17], def[17], att[18], def[18], att[19], def[19], att[20], def[20]
internals = lp, n_steps, is_accept, acceptance_rate, log_density, hamiltonian_energy, hamiltonian_energy_error, max_hamiltonian_energy_error, tree_depth, numerical_error, step_size, nom_step_size

Summary Statistics

parameters		mean	std	mcse	ess_bulk	ess_tail	
rhat	...						
	Symbol	Float64	Float64	Float64	Float64	Float64	Float64
1.0155	...						
	μ att	-0.0106	0.1006	0.0059	286.0900	721.8933	
1.0031	...						
	μ def	-0.0008	0.0994	0.0059	288.4936	464.7421	
1.0023	...						
	σ att	0.3849	0.0779	0.0035	621.2962	339.6839	
1.0028	...						
	σ def	0.2154	0.0612	0.0034	327.5896	117.0460	
	home	0.3380	0.0425	0.0009	2208.4531	2263.3756	

1.0004	...					
	att[1]	-0.2672	0.2014	0.0100	409.4098	955.1314
1.0069	...					
	def[1]	-0.1680	0.1656	0.0079	421.9315	441.1022
1.0004	...					
	att[2]	-0.0720	0.1883	0.0100	355.6188	657.8336
1.0078	...					
	def[2]	0.0594	0.1556	0.0078	400.4448	1113.5101
1.0011	...					
	att[3]	-0.4594	0.2072	0.0103	402.8109	1040.9419
1.0095	...					
	def[3]	0.0862	0.1563	0.0078	406.0630	531.2944
1.0042	...					
	att[4]	0.5008	0.1743	0.0101	295.1097	771.1988
1.0159	...					
	def[4]	-0.1551	0.1678	0.0081	425.6426	571.6784
0.9998	...					
	att[5]	0.2910	0.1777	0.0102	306.1388	709.1792
1.0109	...					
	def[5]	-0.3300	0.1843	0.0084	445.9616	167.5677
0.9999	...					
	att[6]	-0.1918	0.1969	0.0104	360.3907	940.1000
1.0115	...					
	def[6]	-0.0497	0.1541	0.0069	500.0782	991.6641
1.0016	...					
	:	:	:	:	:	:
...						

1 column and 28 rows

s omitted

Quantiles

parameters	2.5%	25.0%	50.0%	75.0%	97.5%
Symbol	Float64	Float64	Float64	Float64	Float64
μatt	-0.2005	-0.0769	-0.0106	0.0569	0.1855
μdef	-0.2022	-0.0655	0.0009	0.0668	0.1866
σatt	0.2633	0.3281	0.3729	0.4302	0.5561
σdef	0.0804	0.1757	0.2123	0.2506	0.3494
home	0.2516	0.3101	0.3398	0.3673	0.4177
att[1]	-0.6797	-0.3978	-0.2636	-0.1333	0.1242
def[1]	-0.4877	-0.2831	-0.1685	-0.0515	0.1470
att[2]	-0.4397	-0.1989	-0.0713	0.0565	0.2913
def[2]	-0.2495	-0.0467	0.0640	0.1688	0.3637
att[3]	-0.8748	-0.5914	-0.4588	-0.3222	-0.0598
def[3]	-0.2396	-0.0133	0.0914	0.1876	0.3825
att[4]	0.1529	0.3881	0.5069	0.6142	0.8376
def[4]	-0.4933	-0.2669	-0.1504	-0.0392	0.1603
att[5]	-0.0623	0.1790	0.2905	0.4048	0.6297
def[5]	-0.6938	-0.4537	-0.3274	-0.2027	0.0213
att[6]	-0.5740	-0.3206	-0.1891	-0.0596	0.1925
def[6]	-0.3658	-0.1516	-0.0457	0.0545	0.2445
:	:	:	:	:	:

28 rows omitted

```
In [20]: posterior_df=DataFrame(posterior)
```

3000×59 DataFrame

2975 rows omitted

Row	iteration	chain	μ att	μ def	σ att	σ def	home
	Int64	Int64	Float64	Float64	Float64	Float64	Float64
1	1001	1	-0.0340967	0.0488872	0.371805	0.341548	0.41
2	1002	1	0.120224	0.0379764	0.442336	0.104175	0.318
3	1003	1	0.120224	0.0379764	0.442336	0.104175	0.318
4	1004	1	0.14776	0.0602034	0.431715	0.100922	0.329
5	1005	1	0.14776	0.0602034	0.431715	0.100922	0.329
6	1006	1	0.14776	0.0602034	0.431715	0.100922	0.329
7	1007	1	0.132189	0.0570596	0.422664	0.106005	0.358
8	1008	1	0.140849	0.0218837	0.340951	0.149594	0.302
9	1009	1	0.000211701	0.00791739	0.55847	0.0972845	0.407
10	1010	1	-0.0436966	-0.0107217	0.500536	0.108407	0.4
11	1011	1	0.198782	0.0145961	0.331272	0.195929	0.260
12	1012	1	-0.00990157	0.133114	0.629202	0.263415	0.343
13	1013	1	0.143073	0.0929881	0.423556	0.258046	0.315
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
2989	3989	1	-0.152038	0.000856577	0.322627	0.227176	0.2
2990	3990	1	0.0130078	0.0471447	0.487366	0.20001	0.379
2991	3991	1	-0.120245	-0.0915178	0.35016	0.21365	0.365
2992	3992	1	-0.0803213	-0.0308569	0.334981	0.158142	0.40
2993	3993	1	-0.0574519	-0.0464465	0.347634	0.225209	0.308
2994	3994	1	0.10894	-0.106624	0.42737	0.23224	0.381
2995	3995	1	0.00348895	-0.0755697	0.39041	0.162871	0.317
2996	3996	1	0.0239133	-0.128041	0.357514	0.194984	0.395
2997	3997	1	0.180213	-0.0651568	0.314539	0.16265	0.37
2998	3998	1	0.0579803	0.0572436	0.451031	0.199167	0.369
2999	3999	1	0.176402	0.0165819	0.444763	0.227071	0.329
3000	4000	1	-0.0406021	-0.0143352	0.464827	0.241888	0.33

In [21]: DataFrames.transform!(posterior_df, AsTable(Between("att[1]", "att[20]"

```
DataFrames.transform!(posterior_df, AsTable(Between("def[1]", "def[20]"  
DataFrames.transform!(posterior_df, AsTable([:att_mean, :def_mean]) =>
```

3000×62 DataFrame

2975 rows omitted

Row	iteration	chain	μ att	μ def	σ att	σ def	home
	Int64	Int64	Float64	Float64	Float64	Float64	Float64
1	1001	1	-0.0340967	0.0488872	0.371805	0.341548	0.41
2	1002	1	0.120224	0.0379764	0.442336	0.104175	0.318
3	1003	1	0.120224	0.0379764	0.442336	0.104175	0.318
4	1004	1	0.14776	0.0602034	0.431715	0.100922	0.329
5	1005	1	0.14776	0.0602034	0.431715	0.100922	0.329
6	1006	1	0.14776	0.0602034	0.431715	0.100922	0.329
7	1007	1	0.132189	0.0570596	0.422664	0.106005	0.358
8	1008	1	0.140849	0.0218837	0.340951	0.149594	0.302
9	1009	1	0.000211701	0.00791739	0.55847	0.0972845	0.407
10	1010	1	-0.0436966	-0.0107217	0.500536	0.108407	0.4
11	1011	1	0.198782	0.0145961	0.331272	0.195929	0.260
12	1012	1	-0.00990157	0.133114	0.629202	0.263415	0.343
13	1013	1	0.143073	0.0929881	0.423556	0.258046	0.315
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
2989	3989	1	-0.152038	0.000856577	0.322627	0.227176	0.2
2990	3990	1	0.0130078	0.0471447	0.487366	0.20001	0.379
2991	3991	1	-0.120245	-0.0915178	0.35016	0.21365	0.365
2992	3992	1	-0.0803213	-0.0308569	0.334981	0.158142	0.40
2993	3993	1	-0.0574519	-0.0464465	0.347634	0.225209	0.308
2994	3994	1	0.10894	-0.106624	0.42737	0.23224	0.381
2995	3995	1	0.00348895	-0.0755697	0.39041	0.162871	0.317
2996	3996	1	0.0239133	-0.128041	0.357514	0.194984	0.395
2997	3997	1	0.180213	-0.0651568	0.314539	0.16265	0.37
2998	3998	1	0.0579803	0.0572436	0.451031	0.199167	0.369
2999	3999	1	0.176402	0.0165819	0.444763	0.227071	0.329
3000	4000	1	-0.0406021	-0.0143352	0.464827	0.241888	0.33

In [22]: `# For this example, we are interested in a pair of teams (no need to u`

```

teamA = "Manchester City"
teamB = "Manchester United"

teamA_id = findfirst(isequal(teamA), teams)
teamB_id = findfirst(isequal(teamB), teams)

teamA_att_post = posterior_df[:, "att[$teamA_id]"]
teamA_def_post = posterior_df[:, "def[$teamA_id]"]

teamB_att_post = posterior_df[:, "att[$teamB_id]"]
teamB_def_post = posterior_df[:, "def[$teamB_id]"]

```

3000-element Vector{Float64}:

```


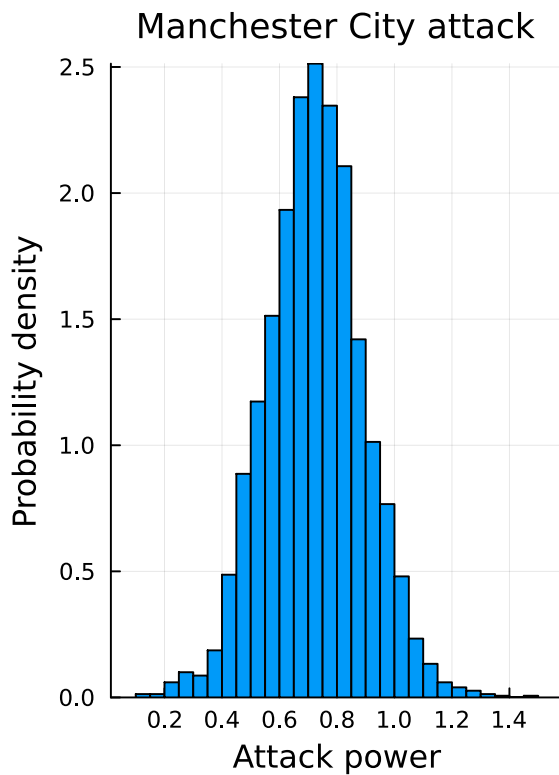
-0.31659686077845584
 0.010909344846748792
 0.010909344846748792
-0.005602923863187977
-0.005602923863187977
-0.005602923863187977
-0.12470192624633054
-0.23946056821044293
-0.17673201401029642
-0.12582298768177858
  ⋮
-0.28907656726617503
-0.600319104560645
-0.25490945685899014
-0.4968341292721567
-0.2274802264310944
-0.7068032570364412
-0.04770233420003743
-0.03970682618439173
-0.5052118038597926

```



```

In [23]: ha1 = histogram(teamA_att_post, title=teamA*" attack", titlefontsize =
ha2 = histogram(teamB_att_post, title=teamB*" attack", titlefontsize =
plot(ha1, ha2, layout=(1,2));
xlabel!("Attack power");
ylabel!("Probability density")

```



```
In [24]: hd1 = histogram(teamA_def_post, title=teamA*" defense", titlefontsize
hd2 = histogram(teamB_def_post, title=teamB*" defense", titlefontsize
plot(hd1, hd2, layout=(1,2));
xlabel!("Defense power");
ylabel!("Probability density")
```



Mini Project

Consult the lecture notes.

```
In [ ]: Random.seed!(205579184)
# hint: let's simulate 500 hypothetical finals (then you will have a t

# first leg: teamA is the home team and teamB is the away team

θ_home = posterior_df[:, :home] + posterior_df[:, "att[$teamA_id]" ] + po
θ_away = posterior_df[:, "att[$teamB_id]" ] + posterior_df[:, "def[$teamA

teamA_score = rand.(Poisson.(exp.(θ_home)), 500)
teamB_score = rand.(Poisson.(exp.(θ_away)), 500)

# second leg: teamA is the away team and teamB is the home team

θ_home = posterior_df[:, :home] + posterior_df[:, "att[$teamB_id]" ] + po
θ_away = posterior_df[:, "att[$teamA_id]" ] + posterior_df[:, "def[$teamB

teamA_score += rand.(Poisson.(exp.(θ_away)), 500) # add the first-leg
teamB_score += rand.(Poisson.(exp.(θ_home)), 500)

# transform into long column vectors
teamA_score = vcat(teamA_score...)
teamB_score = vcat(teamB_score...)


display(histogram2d(teamA_score, teamB_score, title=teamA*" vs "*teamB
# https://docs.juliaplots.org/dev/generated/colorschemes/

# Winning probabilities
winning_prob_A = sum(teamA_score .> teamB_score) / length(teamA_score)
println("Winning probability of "*teamA*" against "*teamB*" is "*string(

winning_prob_B = sum(teamA_score .< teamB_score) / length(teamA_score)
println("Winning probability of "*teamB*" against "*teamA*" is "*string(

draw_prob = sum(teamA_score .== teamB_score) / length(teamA_score)
println("Draw probability between "*teamA*" and "*teamB*" is "*string(

println("Sum of probabilities (sanity check): "*string((winning_prob_A
```



Winning probability of Manchester City against Manchester United is 61.504%

Winning probability of Manchester United against Manchester City is 24.09%

Draw probability between Manchester City and Manchester United is 14.405%

Sum of probabilities (sanity check): 100.0%